

Tejesvani Muppara Vijayaram

Boston, MA(Willing to relocate) | mvtejesvani@gmail.com | (857) 465 9450 | [LinkedIn](#) | [GitHub](#)

SUMMARY

Data Engineer with 3 years of experience building and supporting scalable ETL/ELT pipelines, real-time streaming systems, and cloud-based analytics infrastructure. Hands-on experience with Apache Airflow, Snowflake, MySQL, Kafka, PySpark, and REST APIs in Agile/Scrum environments. Skilled in developing workflow orchestration pipelines, optimizing SQL queries, monitoring streaming data systems, and improving data quality for analytics and operational reporting.

Core stack: SQL, Python, dbt, Airflow, Docker, REST API

WORK EXPERIENCE

myRide

Boston, MA

Data Engineer

Oct 2025 - Present

- Cut pipeline execution time by 42% on 1TB/day sensor and ride data ingestion by refactoring Airflow DAGs to process data as micro-batches in parallel, enabling near real-time data availability for downstream ML use cases.
- Improved performance metric reliability by 32% for engineering teams by building new dbt staging and intermediate models that enforced null checks, deduplication, and referential integrity across ride, vehicle, and pricing sources.

LTIMindtree

Chennai, TN

Data Engineer

Jan 2022 - May 2023

- Built scalable ETL pipelines that ingested data from 6+ source systems spanning policy, claims, and underwriting domains into Snowflake using AWS Glue and Airflow, establishing analytics-ready data for reporting teams.
- Reduced end-to-end pipeline runtime by 57% by implementing incremental CDC ingestion pipelines using AWS DMS and Glue jobs, eliminating unnecessary full reloads and improving data freshness for downstream consumers.
- Optimized Snowflake query runtime by 64% across reporting workloads by redesigning staging transformations with partition-aware SQL models and standardized dimensional structures, improving dashboard responsiveness.
- Increased enterprise data quality from 82% to 94% by implementing automated SQL validation checks, schema consistency testing, and monitoring workflows across critical ETL pipelines.

Associate Engineer

Aug 2021 - Dec 2021

- Streamlined transformation workflows for 10+ insurance datasets by defining relational mappings, SQL transformation rules, and validation logic across ETL pipelines, reducing reporting inconsistencies by 25%.
- Resolved 40+ production data issues through root cause analysis, reconciliation checks, and cross-system SQL validation across S3, Redshift, and reporting layers, reducing reporting delays by approximately 35%.

PROJECTS

Transaction Fraud Detection | *Kafka, PostgreSQL, Docker, Kubernetes, Grafana, Prometheus*

[GitHub](#)

- Built a Kafka and Spark Streaming fraud pipeline processing 30–50 events/min with DLQ and replay handling, achieving 98% processing reliability and eliminating data loss across failure scenarios.
- Improved transaction validation accuracy by 95% by developing Python microservices with reconciliation logic, audit trails, and idempotent PostgreSQL writes for real-time fraud event processing.
- Reduced incident detection time by 70% by implementing Prometheus and Grafana monitoring for Kafka lag, throughput, and pipeline failures, improving operational visibility during traffic spikes.

Clinical Data AWS Platform | *AWS ECS/Fargate, RDS, S3, Terraform, Docker, GitHub Actions*

[GitHub](#)

- Architected a containerized clinical trials data pipeline on AWS ECS Fargate, provisioning 35 cloud resources across VPC, RDS, KMS, Secrets Manager via Terraform, reducing the infrastructure setup to a single command.
- Built modular Python pipeline ingesting 5K clinical trial records from S3 into RDS PostgreSQL with automated transformation, KMS-encrypted secrets, and CloudWatch monitoring, achieving zero-hardcoded credentials.
- Implemented GitHub Actions CI/CD with 21 unit tests on every push and integration tests against live AWS — automating Docker builds, ECR pushes, and ECS deployments on every merge to main.

Bank Marketing Prediction | *GCP, BigQuery, Airflow, Docker*

[GitHub](#)

- Orchestrated an end-to-end ML pipeline on GCP using Airflow DAGs and Docker containers, automating data ingestion, feature engineering, model training, and evaluation with full lineage tracking in BigQuery.
- Redesigned BigQuery transformation logic with partitioned and clustered tables, reducing ETL job runtime by 70% and enabling automated campaign performance reporting without manual intervention.

TECHNICAL SKILLS

Data Engineering & Orchestration: Airflow, dbt, Kafka, Snowflake, Data Models

Cloud Platforms: AWS (Redshift, RDS, Glue, S3, KMS), Azure (Synapse Analytics, ADF, ADLS), GCP (BigQuery)

Programming & Databases: Python (Pandas, NumPy), SQL (PostgreSQL, MySQL, Oracle),

DevOps & Collaboration: Git, GitHub Actions, CI/CD, Docker, Kubernetes, Terraform(IaC), REST API, JIRA, Confluence

Data Visualization and Monitoring: Power BI, Prometheus, Grafana

EDUCATION

Northeastern University, Boston

Master of Science in Data Analytics Engineering

Sept 2023 - May 2025

CGPA 3.95