

Tejesvani Muppara Vijayaram

Boston, MA(Willing to relocate) | mvtejesvani@gmail.com | (857) 465 9450 | [LinkedIn](#) | [GitHub](#)

SUMMARY

Data Engineer with 3+ years of experience building and optimizing production-grade data pipelines, dimensional models, and analytics infrastructure. Experience processing large-scale transactional and sensor datasets with a focus on data quality, schema governance, and analytics-ready modeling. **Core stack:** SQL, Python, dbt, Azure, Kafka, Kubernetes

WORK EXPERIENCE

myRide

Boston, MA

Oct 2025 - Present

Data Engineer

- Cut pipeline execution time by 42% on 1TB/day sensor and ride data ingestion by refactoring Airflow DAGs to process data as micro-batches in parallel, enabling near real-time data availability for downstream ML use cases.
- Improved performance metric reliability by 32% for engineering teams by building new dbt staging and intermediate models that enforced null checks, deduplication, and referential integrity across ride, vehicle, and pricing sources.
- Reduced bad sensor data by 53% by building a dbt validation layer on Snowflake, flagging anomalies, out-of-range values, and late-arriving records before reporting.

LTIMindtree

Chennai, TN

Jan 2022 - May 2023

Data Engineer

- Migrated 50M+ insurance records from on-prem Oracle to ADLS Gen2 using Azure Data Factory with Delta format, replacing slow and unreliable legacy reporting with a lineage-tracked, audit-ready data foundation.
- Cut pipeline runtime by 57% by replacing end-of-day full reloads with incremental CDC using ADF change tracking, eliminating the processing overhead of reloading unchanged records on every run.
- Standardized metric definitions used by finance and claims teams by building 30+ dbt models on Synapse across Silver and Gold layers, including SCD Type 2 logic for policy dimension tables to preserve full historical changes.
- Raised data quality from 82% to 94% across the Medallion pipeline by implementing dbt tests at every layer and ADF failure alerts, cutting the manual reporting corrections the team had to make every week.
- Reduced deployment time by 50% by provisioning Azure infrastructure with Terraform and automating dbt deployments through Azure DevOps CI/CD, making deployments consistent and repeatable.

Associate Engineer

Aug 2021 - Dec 2021

- Designed source-to-target mappings and logical data models for 10+ datasets by defining transformation rules and business logic, enabling accurate pipeline design and reducing downstream rework by ~25%
- Monitored ADF ETL pipelines for underwriting and claims datasets by validating scheduled run outputs, triaging failures, and escalating schema issues — building deep familiarity with the end-to-end pipeline architecture.
- Resolved 40+ data availability and consistency issues by tracing discrepancies across source and target systems using SQL and validation checks, reducing reporting delays by ~35%.

PROJECTS

Transaction Fraud Detection | Kafka, PostgreSQL, Docker, Kubernetes, Grafana, Prometheus

[GitHub](#)

- Architected distributed Kafka streaming pipeline with Avro schema governance, fault-tolerant DLQ/offset replay, and idempotent PostgreSQL writes, achieving 98% processing reliability with zero data loss across failure scenarios.
- Built Python risk-scoring microservices processing 30–50 events/min with full audit trails and reconciliation logic across Kafka and PostgreSQL, enforcing end-to-end data integrity, reducing undetected inconsistencies by 95%.
- Scaled pipeline to handle 2–3x traffic spikes using Kubernetes HPA and reduced incident detection time by 70% via Prometheus + Grafana dashboards monitoring consumer lag, throughput, and error rates.

Bank Marketing Prediction | GCP, BigQuery, Airflow, Docker

[GitHub](#)

- Orchestrated an end-to-end ML pipeline on GCP using Airflow DAGs and Docker containers, automating data ingestion, feature engineering, model training, and evaluation with full lineage tracking in BigQuery.
- Redesigned BigQuery transformation logic with partitioned and clustered tables, reducing ETL job runtime by 70% and enabling automated campaign performance reporting without manual intervention.

TECHNICAL SKILLS

Data Engineering & Orchestration: Airflow, dbt, PySpark, Kafka, Snowflake, Delta Lake, Data Models

Cloud Platforms: Azure (Synapse Analytics, ADF, ADLS, Databricks), GCP (BigQuery), AWS (Redshift)

Programming & Databases: Python (Pandas, NumPy), SQL (PostgreSQL, MySQL, Oracle), R, Scala

DevOps & Collaboration: Git, GitHub Actions, CI/CD, Docker, Kubernetes, Terraform(IaC), REST API, JIRA, Confluence

Data Visualization and Monitoring: Power BI (DAX, Power Query), Tableau, Prometheus, Grafana

AI Tools: Claude Code, Claude SKILLS, OpenAI, Gemini

EDUCATION

Northeastern University, Boston

Sept 2023 - May 2025

Master of Science in Data Analytics Engineering

CGPA 3.95

Anna University, Chennai, TN, India

Sept 2017 - May 2021

Bachelors of Engineer in Computer Science and Engineering

CGPA 3.64